

LITO ZARATE BIALA

AI Engineer | GenAI in Production | Enterprise SRE Foundation

Minneapolis–St. Paul, MN · lzbiala@gmail.com · linkedin.com/in/lzbiala · github.com/LZBiala · lzbiala.github.io

SUMMARY

AI engineer who ships GenAI to production inside a regulated enterprise: at Wells Fargo I designed and deployed GenAI-powered incident root-cause analysis with human-in-the-loop validation, running against production incidents — backed by 12+ years of leadership and 5+ years of SRE and platform practice. I also build AI systems that publish their own measurements: four open-source, CI-verified LLM evaluation and memory harnesses that reproduce every published result and fail the build when a claim drifts.

EXPERIENCE

Wells Fargo — Lead Technology Business Systems Consultant

SRE & incident management leadership · Minneapolis, MN | 2021 – Present

- **Production GenAI incident RCA:** designed and deployed a human-in-the-loop system that generates root-cause analysis from production incidents, supports validation before conclusions are used, and speeds incident-leadership reviews.
- **Enterprise AI adoption:** led early RAG adoption and integration with enterprise systems through risk and compliance gates; led agentic-AI automation for compliance workflows to improve audit readiness.
- **Reliability foundation:** led incident management and 5-why root-cause analysis for customer-critical applications, reducing incident response times (MTTx); maintained log-based alerting and monitoring through platform migrations to preserve operational continuity.
- **Platform engineering:** engineered foundational platforms for scale through cloud migration and data-center exit; led engineering release meetings across AppDev, Platform, and Product teams toward production readiness via CI/CD.
- **Cross-functional leadership:** guided distributed US/India engineering teams and translated incident, reliability, and AI-system findings into decision-ready reviews for engineering and executive stakeholders.

Earlier career (pre-2021) — leadership & team management roles

People and team leadership; full history on LinkedIn.

OPEN-SOURCE AI PROJECTS — GITHUB.COM/LZBIALA (ALL LIVE, CI-VERIFIED)

Four harnesses, one architecture: memory → evaluation → decision calibration → the control loop. Every published number regenerates in CI with zero API keys; each repo ships an interactive walkthrough a stakeholder can click through.

- **wiki-memory-lab:** agent long-term memory as an auditable markdown wiki — selective recall at 0.61× the context tokens of loading everything, with the crossover where the design *loses* (1.27×) published in the same table.
- **memory-repair-lab:** the model never learns; the notes do — retrieval failures drive cited repairs, re-measured on a held-out split sealed before generation zero: 8/10 → 10/10, placebo-controlled.
- **agent-mutation-lab:** mutation testing where the subject is an AI code reviewer — 9 documented defect classes, byte-identical clean controls, 0 false alarms; one deliberate blind spot so the scorecard shows real failure output.
- **adversarial-chambers:** multi-agent orchestration with the judge on trial — role separation, a kill ledger with mandatory written reasons, and an instruction-flip test that detects instruction-following judges (0/8 vs 3/8 contrast).

SKILLS

AI agent evaluation · LLM observability · evals & benchmark design · RAG & agent memory · multi-agent orchestration · prompt & judge calibration · GenAI incident RCA · Python · pytest · CI/CD (GitHub Actions) · SRE & monitoring (Splunk, Grafana) · SLOs, MTtx & incident management · ITIL · platform engineering & cloud migration · distributed team leadership · executive communication

EDUCATION & CREDENTIALS

Data Analytics Boot Camp, University of Minnesota · B.A. Psychology, Metropolitan State University

Splunk certified (Fundamentals 1–3, Cloud Observability) · Wells Fargo internal SRE-leadership and AI Dojo engineering pathways