

LITO ZARATE BIALA

HANDS-ON AI BUILDER & ENTERPRISE AI PRODUCT LEAD | LLM EVALUATION FRAMEWORKS · AGENTS · RAG · MCP · SRE RELIABILITY

Minneapolis–St. Paul, MN · lzbiala@gmail.com · linkedin.com/in/lzbiala · github.com/LZBiala · lzbiala.github.io

SUMMARY

Hands-on AI builder and enterprise AI product lead at Wells Fargo: I ship governed GenAI products end to end — problem framing, workflow design, hands-on implementation, evaluation suites, and stakeholder adoption — backed by 12+ years of leadership across people management and technical delivery. Reliability-first AI grounded in human review, evidence provenance, observability, and enterprise risk. Public, CI-verified evaluation portfolio at github.com/LZBiala — the build fails if a published claim drifts.

AI LEADERSHIP & TECHNICAL STRENGTHS

- **AI Product & Strategy:** use-case discovery, roadmap sequencing, requirements, acceptance criteria, adoption, stakeholder alignment.
- **GenAI & Agentic Systems:** RAG, LLM evaluation frameworks and evals, prompt/context engineering, MCP integration, agents, human-in-the-loop controls.
- **AI Engineering Delivery:** Python, Streamlit, TypeScript, GitHub, CI/CD, automated tests, code review, iterative PR-based delivery.
- **Reliability & Governance:** SRE, incident management, 5-Why/RCA, MTTx, observability, fail-closed behavior, AI risk and compliance readiness.
- **Agentic Build Tooling:** Claude Code, Codex, Devin — daily hands-on agentic system delivery with human-in-the-loop controls.

PROFESSIONAL EXPERIENCE

Vice President, Lead Digital Product Service Manager | Wells Fargo, Digital Technology & Innovation

Individual-contributor leadership role · Minneapolis, MN | Feb 2025 – Present

- **AI product delivery:** lead delivery of two governed AI products — an incident-notification quality assistant and a Daily 5-Why QA platform — turning operational pain points into buildable specs: requirements, evidence-source rules, versioned scoring rubrics, and acceptance tests, all gated by human review.
- **Hands-on build:** daily engineering with GitHub, Jira, Devin, VS Code Copilot, Python, Streamlit, and TypeScript — focused branches, pull requests, automated tests, coverage, linting, type checks, and protected-file audits for reviewable change.
- **Evidence architecture:** designed the evidence architecture pulling live incident and RCA data from ServiceNow and SharePoint: every record carries its provenance and passes schema validation; the system redacts diagnostics, fails closed on bad data, and requires explicit approval gates before touching production.
- **Responsible-AI boundaries:** established product boundaries that removed person-level scoring and identity from published score-bearing outputs — keeping names only for routing work to the right team and supporting 24-hour handoffs across global time zones.
- **Incident leadership:** lead major-incident, vulnerability-remediation, and operating-review readiness across distributed stakeholders; translate technical evidence into decision-ready impact, root-cause, detection-gap, mitigation, and preventive-action narratives.
- **AI enablement:** advance AI adoption through executive communication, pair programming, prompt-engineering enablement, technical documentation, and coaching on responsible human-in-the-loop delivery.

Senior Assistant Vice President, Digital Product Service Manager | Wells Fargo

Minneapolis, MN | Oct 2022 – Feb 2025

- **Reliability:** improved resiliency and uptime for customer-critical applications; reduced incident response times (MTTx) through proactive engineering and 5-why root-cause discipline.
- **Observability:** enhanced log-based alerting and monitoring through platform migrations (Splunk, Grafana), preserving operational continuity with executive-level service-health insight.
- **Platform delivery:** engineered foundational platforms for scale through cloud migration and data-center exit; led engineering release meetings across AppDev, Platform, and Product teams via CI/CD; guided distributed US/India engineering teams.

Senior Technology Business Systems Consultant (contract) | Wells Fargo

Jan 2022 – Oct 2023

- Kept digital information systems running smoothly under SRE, ITIL, and ITSM best practices; supported incident and problem management for production applications.

Software Product Analyst (contract)

Mar 2021 – Jan 2022

- Monitored health and status of internal APIs through dashboards and data solutions; analyzed historical trends to surface persistent performance issues.

Earlier people-leadership and operations roles — full history on LinkedIn.

SELECTED ENTERPRISE AI INITIATIVES

AI-Assisted Incident Management Workflows

Reliability, remediation & API gateway lane

- Apply AI workflows across the incident lifecycle — triage, root-cause analysis support, and remediation follow-through — reducing incident response times (MTTx) and improving uptime for critical systems.
- Years managing an API gateway platform made me the connective layer between Product, Application Development, DevOps, incident management, and governance — translating each domain's constraints into decisions the others can act on.

Incident Quality Assurance Platform

Daily 5-Why QA product lane

- Defined phased delivery: a build plan mapping features to phases, version-controlled scoring rubrics, rules requiring every root-cause claim to cite its source evidence, quality checks on preventive actions, and acceptance testing — outputs flagged as advisory pending human review.
- Encoded evidence-aware states so incomplete or non-evaluable content is never misrepresented as passing or failing, improving trust and auditability.

AI Enablement & Architecture Intelligence

Enterprise learning & solution design

- Created reusable prompt-engineering guidance and supported team AI learning; explored enterprise RAG, chunking, metadata enrichment, versioning, and quality checks for architecture-intelligence use cases.

PUBLIC AI EVALUATION & RELIABILITY PORTFOLIO — [GITHUB.COM/LZBIALA](https://github.com/lzbiala) · [LZBIALA.GITHUB.IO](https://lzbiala.github.io)

- wiki-memory-lab**: agent-memory evaluation framework; selective recall at 0.61× the tokens of the full-context baseline, with the loss trade-off (1.27×) measured and published rather than hidden.
- agent-mutation-lab**: behavioral defect-detection harness; caught 9/10 planted defects with 0 false alarms on byte-identical clean controls; best-of-k voting study where panel accuracy tracks the exact binomial (Condorcet) prediction, with published controls showing when voting cannot help.
- adversarial-chambers**: adversarial robustness evals; instruction-flip test — a scripted evidence-driven judge flips 0/8 verdicts, a scripted deferential judge flips 3/8, exactly the ambiguous items.
- memory-repair-lab**: placebo-controlled self-repair evaluation; 8/10 → 10/10 on a held-out split sealed before generation zero.

All four are CI-verified on Windows and Linux: the build fails if any published claim drifts from regenerated data. Each ships an interactive walkthrough a stakeholder can click through.

TECHNOLOGY

AI/ML: Generative AI, RAG, agentic automation, LLM evaluation frameworks, prompt optimization, responsible AI · **Engineering**: Python, Streamlit, TypeScript, Git, GitHub, APIs, CI/CD, automated testing · **Operations**: Splunk, Grafana, observability, incident/problem management, 5-Why RCA, MTtx, ITIL · **Leadership**: AI product strategy, distributed-team leadership, technical writing, executive communication, risk and compliance · **Agentic tooling**: Claude Code, Codex, Devin (Cognition)

EDUCATION & CREDENTIALS

Data Analytics & Data Visualization Boot Camp Certificate, University of Minnesota Twin Cities · B.A. Psychology, Metropolitan State University

Splunk courses: Fundamentals 1–3 & Cloud Observability · Anthropic Academy course completions, 2026: Claude Code in Action · Introduction to Model Context Protocol · Wells Fargo internal SRE-leadership and AI-engineering program pathways