

LITO ZARATE BIALA

HANDS-ON AI BUILDER & SITE RELIABILITY ENGINEERING LEAD | AGENTIC SYSTEMS · MCP · INCIDENT MANAGEMENT · ITIL/ITSM · FORWARD DEPLOYED ENGINEERING

Minneapolis-St. Paul, MN · lzbiala@gmail.com · linkedin.com/in/lzbiala · github.com/LZBiala · lzbiala.github.io

SUMMARY

Hands-on AI builder and site reliability lead at Wells Fargo. I find the bottleneck that needs a foundational fix, build the AI that clears it, and make the system prove it worked. Day to day I run SRE and incident practice as the working link between Product, Platform, Development, and Infrastructure: I turn each group's constraints into decisions the others can act on, and I ship governed GenAI products end to end (problem framing, workflow design, hands-on implementation, evaluation suites, MTTx metrics, stakeholder adoption), backed by 12+ years of leadership across people management and technical delivery. Reliability-first AI grounded in human review, evidence provenance, observability, ITIL/ITSM, and enterprise risk. Public, CI-verified evaluation portfolio at github.com/LZBiala - the build fails if a measured claim drifts.

AI LEADERSHIP & TECHNICAL STRENGTHS

- **AI Product & Strategy:** use-case discovery with the teams who own the problem, requirements, acceptance criteria, roadmap sequencing, adoption, executive and non-engineer communication.
- **GenAI & Agentic Systems:** RAG, LLM evaluation frameworks and evals, prompt/context engineering, MCP integration, agents, human-in-the-loop controls.
- **AI Engineering Delivery:** Python, Streamlit, TypeScript, GitHub, CI/CD, automated tests, code review, iterative PR-based delivery.
- **Reliability & Governance:** SRE, DevOps, incident management, 5-Why/RCA, MTTx, observability, fail-closed behavior, AI risk and compliance readiness.
- **Agentic Build Tooling:** Claude Code, Codex, Devin - daily hands-on agentic system delivery with human-in-the-loop controls.

PUBLIC AI EVALUATION & RELIABILITY PORTFOLIO - [GITHUB.COM/LZBIALA](https://github.com/LZBiala) · [LZBIALA.GITHUB.IO](https://lzbiala.github.io)

- **wiki-memory-lab:** agent-memory evaluation framework; selective recall at 0.61× the tokens of the full-context baseline, with the loss trade-off (1.27×) measured and published rather than hidden.
- **agent-mutation-lab:** behavioral defect-detection harness; caught 9/10 planted defects at the planted line (within two lines), with 0 false alarms on byte-identical clean controls; a separate best-of-k study shows nine independently-noised copies climb exactly as the jury arithmetic predicts while a copycat control stays flat.
- **adversarial-chambers:** proposer, skeptic, and judge with no role played twice; instruction-flip test on two scripted judges, true by construction: an evidence-led judge flips 0/8 verdicts, a deferential judge flips 3/8, exactly the ambiguous items. Proves the instrument detects instruction sensitivity, not that any live judge passes.
- **memory-repair-lab:** placebo-controlled self-repair evaluation; 8/10 to 10/10 on a held-out split sealed before generation zero.

All four are CI-verified on Windows and Linux (the build fails if any measured claim drifts from regenerated data); each ships a walkthrough a stakeholder can click through.

PROFESSIONAL EXPERIENCE

Vice President, Lead Digital Product Service Manager | Wells Fargo, Digital Technology & Innovation

Individual-contributor leadership role · Minneapolis, MN | Feb 2025 - Present

- **AI product delivery:** lead delivery of two governed AI products, an incident-notification quality assistant and a Daily 5-Why QA platform, turning operational bottlenecks into buildable specs gated by human review and grounded in SRE and DevOps practice.
- **Hands-on build:** daily engineering with GitHub, Jira, Devin, VS Code Copilot, Python, Codex, and Claude Code - focused branches, pull requests, automated tests, coverage, linting, type checks, and protected-file audits for reviewable change.
- **Evidence architecture:** designed the evidence architecture pulling live incident and RCA data from ServiceNow and SharePoint: every record carries its provenance and passes schema validation; the system redacts diagnostics, fails closed on bad data, and requires explicit approval gates before touching production.
- **Responsible-AI boundaries:** established product boundaries that removed person-level scoring and identity from published score-bearing outputs, keeping names only for routing work to the right team.
- **Incident leadership:** lead major-incident, vulnerability-remediation, and operating-review readiness across distributed stakeholders; translate technical evidence into decision-ready impact, root-cause, detection-gap, mitigation, and preventive-action narratives. Years managing an API gateway platform made me the connective layer between Product, AppDev, DevOps, and governance.

Senior Assistant Vice President, Digital Product Service Manager | Wells Fargo

Minneapolis, MN | Oct 2022 - Feb 2025

- **Reliability:** improved resiliency and uptime for customer-critical applications; reduced incident response times (MTTx) through proactive engineering, 5-why root-cause discipline, and DevOps practice.
- **Observability:** enhanced log-based alerting and monitoring through platform migrations (Splunk, Grafana), preserving operational continuity with executive-level service-health insight.
- **Platform delivery:** engineered foundational platforms for scale through cloud migration and data-center exit; led engineering release meetings across AppDev, Platform, and Product teams via CI/CD; guided distributed US/India engineering teams.

Senior Technology Business Systems Consultant, Wells Fargo, Digital Technology & Innovation (Jan 2022 - Oct 2022): incident and problem management for production applications under SRE, ITIL, and ITSM practice. **Software Product Analyst**, Wells Fargo, Strategy, Digital & Innovation (Mar 2021 - Jan 2022): monitored internal API health through dashboards and data solutions; surfaced persistent performance issues from historical trends. Earlier people-leadership and operations roles - full history on LinkedIn.

SELECTED ENTERPRISE AI INITIATIVES (2026)

Incident Quality Assurance Platform | Daily 5-Why QA product lane, Wells Fargo

- **Phased delivery:** defined the phased build plan: features mapped to phases, version-controlled scoring rubrics, every root-cause claim required to cite its source evidence, quality checks on preventive actions, and acceptance testing; outputs ship as advisory pending human review.
- **Evidence-aware states:** encoded states so incomplete or non-evaluable content is never misrepresented as passing or failing, improving trust and auditability.

AI Enablement & Architecture Intelligence | Enterprise learning and solution design

- **Enablement:** wrote reusable prompt-engineering guidance and coached engineers on responsible human-in-the-loop delivery through pair programming and executive communication.
- **Solution design:** explored enterprise RAG for architecture-intelligence use cases: chunking, metadata enrichment, versioning, quality checks, and deep-wiki agentic memory.

TECHNOLOGY

AI/ML: Generative AI, RAG, agentic automation, LLM evaluation frameworks, prompt optimization, responsible AI, MCP, knowledge graphs, deep-wiki agent long-term memory · **Engineering:** Python, Streamlit, TypeScript, Git, GitHub, GitHub Actions, APIs, CI/CD, automated testing · **Operations:** SRE, DevOps, Splunk, Grafana, observability, incident/problem management, 5-Why RCA, MTTx, ITIL/ITSM · **Leadership:** AI product strategy, distributed-team leadership, technical writing, executive communication, risk and compliance · **Agentic tooling:** Claude Code, Codex, Devin (Cognition)

EDUCATION & CREDENTIALS

Data Analytics & Data Visualization Boot Camp Certificate, University of Minnesota Twin Cities · B.A. Psychology, Metropolitan State University

Splunk courses: Fundamentals 1-3 & Cloud Observability · Anthropic Academy course completions, 2026: Claude Code in Action · Introduction to Model Context Protocol · Wells Fargo internal SRE-leadership and AI-engineering program pathways