

LITO ZARATE BIALA

HANDS-ON AI BUILDER & SITE RELIABILITY ENGINEERING LEAD | AGENTIC SYSTEMS · MCP · INCIDENT MANAGEMENT · ITIL/ITSM · FORWARD DEPLOYED ENGINEERING

Minneapolis-St. Paul, MN · lzbiala@gmail.com · linkedin.com/in/lzbiala · github.com/LZBiala · lzbiala.github.io

SUMMARY

Hands-on AI builder and site reliability lead at Wells Fargo. I find the bottleneck that needs a foundational fix, build the AI that clears it, and make the system prove it worked. Day to day I run SRE and incident practice as the working link between Product, Platform, Development, and Infrastructure: I turn each group's constraints into decisions the others can act on, and I ship governed GenAI products end to end, backed by 12+ years of leadership across people management and technical delivery. The same approach took major incidents on the API gateway platform I own from 10 in 2024 to 1 year to date in 2026, and cut mean time to detect and react by 40%. Public, CI-verified evaluation portfolio at github.com/LZBiala - the build fails if a measured claim drifts.

AI LEADERSHIP & TECHNICAL STRENGTHS

AI Product	Use-case discovery with the teams who own the problem, requirements, acceptance criteria, roadmap sequencing, adoption, executive and non-engineer communication.
Agentic AI	RAG, LLM evaluation frameworks and evals, prompt and context engineering, MCP integration, agents, human-in-the-loop controls.
Delivery	Python, Streamlit, TypeScript, GitHub, CI/CD, automated tests, code review, iterative PR-based delivery.
Reliability	SRE, DevOps, incident management, 5-Why/RCA, MTTx, observability, fail-closed behavior, AI risk and compliance readiness.
Tooling	Claude Code, Codex, Devin - daily hands-on agentic system delivery with human-in-the-loop controls.

PUBLIC AI EVALUATION & RELIABILITY PORTFOLIO · [GITHUB.COM/LZBIALA](https://github.com/LZBiala) · [LZBIALA.GITHUB.IO](https://lzbiala.github.io)

- wiki-memory-lab:** agent-memory evaluation framework; selective recall at 0.61× the tokens of the full-context baseline, with the loss trade-off (1.27×) measured and published rather than hidden.
- agent-mutation-lab:** behavioral defect-detection harness; caught 9/10 planted defects at the planted line (within two lines), with 0 false alarms on byte-identical clean controls; a separate best-of-k study shows nine independently-noised copies climb exactly as the jury arithmetic predicts while a copycat control stays flat.
- adversarial-chambers:** proposer, skeptic, and judge with no role played twice; instruction-flip test on two scripted judges, true by construction: an evidence-led judge flips 0/8 verdicts, a deferential judge flips 3/8, exactly the ambiguous items. Proves the instrument detects instruction sensitivity, not that any live judge passes.
- memory-repair-lab:** placebo-controlled self-repair evaluation; 8/10 to 10/10 on a held-out split sealed before generation zero.

All four are CI-verified on Windows and Linux: the build fails if a measured claim drifts from regenerated data. Each ships a walkthrough a stakeholder can click through.

PROFESSIONAL EXPERIENCE

Vice President, Lead Digital Platform Services Manager

Feb 2025 - Present

Wells Fargo, Digital Technology & Innovation · Individual-contributor leadership role · Minneapolis, MN

- AI product delivery:** lead delivery of two governed AI products, an incident-notification quality assistant and a Daily 5-Why QA platform, turning operational bottlenecks into buildable specs gated by human review and grounded in SRE and DevOps practice.
- Hands-on build:** daily engineering in Python with Claude Code, Codex, and Devin: focused branches, pull requests, tests, coverage, and type checks, so every change arrives reviewable.
- Governance by design:** built the evidence architecture behind both products, pulling live incident and RCA data from ServiceNow and SharePoint so every record carries its provenance and passes schema validation, redacts diagnostics, fails closed on bad data, and needs explicit approval before touching production. I also drew the boundary that kept person-level scoring out of published outputs, keeping names only for routing work.
- On-call AI workflows:** built AI workflows that tell an on-call engineer who to escalate to and when, measured from impact start rather than from ticket creation, cutting mean time to mitigate and mean time to react for me and my teammates during major incidents.

- **Incident reduction:** major incidents escalated against the API gateway platform I own fell from 10 in 2024 to 4 in 2025 to 1 year to date in 2026, on one of the noisiest and most business-critical surfaces in the enterprise.
- **Escalation design:** lead major-incident, vulnerability-remediation, and operating-review readiness; wrote escalation playbooks an on-call engineer can pick up cold and use to reach the right team in minutes, built on a working map of the organization's critical applications and their dependencies. I carry weekend on-call for those systems.
- **Executive reporting:** run 5-why root-cause analysis to the true cause, then present findings, preventive actions, and impact trends to senior leadership so remediation investment follows the evidence rather than the loudest incident.

Senior Assistant Vice President, Digital Product Service Manager

Oct 2022 - Feb 2025

Wells Fargo · Minneapolis, MN

- **Reliability:** cut mean time to detect and mean time to react by 40% on customer-critical applications through proactive engineering, 5-why root-cause discipline, and DevOps practice.
- **Observability:** enhanced log-based alerting and monitoring through platform migrations (Splunk, Grafana, AppDynamics), preserving operational continuity with executive-level service-health insight.
- **Platform delivery:** engineered foundational platforms for scale through cloud migration and data-center exit; led engineering release meetings across AppDev, Platform, and Product teams via CI/CD; guided distributed US/India engineering teams.

Senior Technology Business Systems Consultant

Jan 2022 - Oct 2022

Wells Fargo, Digital Technology & Innovation

Incident and problem management for production applications under SRE, ITIL, and ITSM practice.

Software Product Analyst

Mar 2021 - Jan 2022

Wells Fargo, Strategy, Digital & Innovation

Monitored internal API health through dashboards and data solutions; surfaced persistent performance issues from historical trends.

Earlier people-leadership and operations roles - full history on LinkedIn.

SELECTED ENTERPRISE AI INITIATIVES (2026)

Incident Quality Assurance Platform

Wells Fargo

Daily 5-Why QA product lane

- **Phased delivery:** defined the phased build plan: features mapped to phases, version-controlled scoring rubrics, every root-cause claim required to cite its source evidence, quality checks on preventive actions, and acceptance testing; outputs ship as advisory pending human review.
- **Evidence-aware states:** encoded states so incomplete or non-evaluable content is never misrepresented as passing or failing, improving trust and auditability.

AI Enablement & Architecture Intelligence

Wells Fargo

Enterprise learning and solution design

- **Enablement:** wrote reusable prompt-engineering guidance and coached engineers on responsible human-in-the-loop delivery through pair programming and executive communication.

TECHNOLOGY

AI/ML	Generative AI, RAG, agentic automation, LLM evaluation frameworks, prompt optimization, responsible AI, MCP, knowledge graphs, deep-wiki agent long-term memory
Engineering	Python, Streamlit, TypeScript, Git, GitHub, GitHub Actions, APIs, CI/CD, automated testing
Operations	SRE, DevOps, Splunk, Grafana, AppDynamics, production triage, observability, incident and problem management, 5-Why RCA, MTTx, ITIL/ITSM
Leadership	AI product strategy, distributed-team leadership, technical writing, executive communication, risk and compliance
Agentic	Claude Code, Codex, Devin (Cognition)

EDUCATION & CREDENTIALS

Data Analytics & Data Visualization Boot Camp Certificate, University of Minnesota Twin Cities · B.A. Psychology, Metropolitan State University

Splunk Certified: Fundamentals 1-3 and Cloud Observability · Wells Fargo internal SRE-leadership and AI-engineering program pathways

Anthropic Academy, 2026: Claude Code in Action · Introduction to Model Context Protocol